

Predict Stock Market with ML

January 2, 2025

1 Predict The Stock Market With ML

We'll going to use S&P 500 and predict tomorrow's index price using 20 years of historical data. We will also backtesting if this is accurate prediction. After that we'll improve the model by adding predictors.

```
[1]: pip install yfinance
```

```
Requirement already satisfied: yfinance in /opt/conda/lib/python3.11/site-  
packages (0.2.51)  
Requirement already satisfied: pandas>=1.3.0 in /opt/conda/lib/python3.11/site-  
packages (from yfinance) (2.1.3)  
Requirement already satisfied: numpy>=1.16.5 in /opt/conda/lib/python3.11/site-  
packages (from yfinance) (1.24.4)  
Requirement already satisfied: requests>=2.31 in /opt/conda/lib/python3.11/site-  
packages (from yfinance) (2.31.0)  
Requirement already satisfied: multitasking>=0.0.7 in  
/opt/conda/lib/python3.11/site-packages (from yfinance) (0.0.11)  
Requirement already satisfied: lxml>=4.9.1 in /opt/conda/lib/python3.11/site-  
packages (from yfinance) (5.2.2)  
Requirement already satisfied: platformdirs>=2.0.0 in  
/opt/conda/lib/python3.11/site-packages (from yfinance) (4.1.0)  
Requirement already satisfied: pytz>=2022.5 in /opt/conda/lib/python3.11/site-  
packages (from yfinance) (2023.3.post1)  
Requirement already satisfied: frozendict>=2.3.4 in  
/opt/conda/lib/python3.11/site-packages (from yfinance) (2.4.6)  
Requirement already satisfied: peewee>=3.16.2 in /opt/conda/lib/python3.11/site-  
packages (from yfinance) (3.17.8)  
Requirement already satisfied: beautifulsoup4>=4.11.1 in  
/opt/conda/lib/python3.11/site-packages (from yfinance) (4.12.2)  
Requirement already satisfied: html5lib>=1.1 in /opt/conda/lib/python3.11/site-  
packages (from yfinance) (1.1)  
Requirement already satisfied: soupsieve>1.2 in /opt/conda/lib/python3.11/site-  
packages (from beautifulsoup4>=4.11.1->yfinance) (2.5)  
Requirement already satisfied: six>=1.9 in /opt/conda/lib/python3.11/site-  
packages (from html5lib>=1.1->yfinance) (1.16.0)  
Requirement already satisfied: webencodings in /opt/conda/lib/python3.11/site-  
packages (from html5lib>=1.1->yfinance) (0.5.1)
```

Requirement already satisfied: python-dateutil>=2.8.2 in
/opt/conda/lib/python3.11/site-packages (from pandas>=1.3.0->yfinance) (2.8.2)
Requirement already satisfied: tzdata>=2022.1 in /opt/conda/lib/python3.11/site-
packages (from pandas>=1.3.0->yfinance) (2023.3)
Requirement already satisfied: charset-normalizer<4,>=2 in
/opt/conda/lib/python3.11/site-packages (from requests>=2.31->yfinance) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /opt/conda/lib/python3.11/site-
packages (from requests>=2.31->yfinance) (3.6)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/opt/conda/lib/python3.11/site-packages (from requests>=2.31->yfinance) (2.1.0)
Requirement already satisfied: certifi>=2017.4.17 in
/opt/conda/lib/python3.11/site-packages (from requests>=2.31->yfinance)
(2023.11.17)
Note: you may need to restart the kernel to use updated packages.

```
[2]: # Import
import yfinance as yf
import pandas as pd
```

```
[3]: sp500 = yf.Ticker("^GSPC") # Download the Price History from Ticker class
sp500 = sp500.history(period = "max")
```

```
[4]: #sp500.index
```

```
[5]: sp500.plot.line(y="Close", use_index = True)
```

```
[5]: <Axes: xlabel='Date'>
```



```

model = RandomForestClassifier(n_estimators = 200, min_samples_split = 100,
    ↪random_state = 1)

# We can't use cross validation. Why? Time series nature of data (train works,
    ↪well but real world would be horrible)
train = sp500.iloc[:-100]
test = sp500.iloc[-100:]

# We aren't going to use Tomorrow or Target since if we use that will cause
    ↪overfit since it's already know the results
predictors = ["Close", "Volume", "Open", "High", "Low"]

model.fit(train[predictors], train["Target"])

```

```
[12]: RandomForestClassifier(min_samples_split=100, n_estimators=200, random_state=1)
```

```

[13]: # Verify the precision. How precise that we predict goes up actually stock
    ↪price go up.
from sklearn.metrics import precision_score

preds = model.predict(test[predictors]) #array

```

```

[14]: preds = pd.Series(preds, index = test.index)

precision_score(test["Target"],preds)

```

```
[14]: 0.7
```

We got precision as 70% which is good but we could make it better


```

[15]: # Plot prediction
combined = pd.concat([test["Target"], preds], axis = 1) #axis = 1 so each test
    ↪output will be as a column

combined.plot()

```

```
[15]: <Axes: xlabel='Date'>
```



1.0.2 Backtesting

```
[16]: # Similar as above but instead using train predictor as well.
def predict(train, test, predictors, model):
    model.fit(train[predictors], train["Target"])
    preds = model.predict(test[predictors])
    preds = pd.Series(preds, index=test.index, name="Predictions") #Named the
    ↪series as Predictions
    combined = pd.concat([test["Target"], preds], axis = 1)
    return combined
```

```
[17]: # Training

#Training for 10 years where each business day is 250 days per year. 25000 days
    ↪for 10 years
def backtest(data, model, predictors, start = 2500, step = 250):
    all_predictions = []

    for i in range(start, data.shape[0], step): #data.shape[0] -> # of rows
        train = data.iloc[0:i].copy()
        test = data.iloc[i:(i+step)].copy()
```

```

        predictions = predict(train, test, predictors, model)
        all_predictions.append(predictions)

    return pd.concat(all_predictions)

```

```

[18]: #Call Back test
      predictions = backtest(sp500, model, predictors)

      predictions["Predictions"].value_counts()

```

```

[18]: Predictions
      0    3628
      1    2689
      Name: count, dtype: int64

```

```

[19]: precision_score(predictions["Target"], predictions["Predictions"])

```

```

[19]: 0.5273335812569728

```

1.0.3 Benchmark

```

[20]: predictions["Target"].value_counts() / predictions.shape[0]

```

```

[20]: Target
      1    0.535856
      0    0.464144
      Name: count, dtype: float64

```

Based on the Benchmark, 52.8% precision make it little worse than natural percentage of the days that stockmarket goes up.

1.0.4 Adding Additional Predictors

By adding more predictors check that whether the accuracy goes up or not.

```

[21]: horizons = [2, 5, 60, 250, 1000] #last 2 days, 5 days, 3 months, etc.

      #By checking for each span ratio, how long it goes up in that spans to help
      ↪better prediction (hypothesis).
      new_predictors = []

      for horizon in horizons:
          rolling_avg = sp500.rolling(horizon).mean() # rolling() provide window
          ↪(days span) calculation

          ratio_column = f"Close_Ratio_{horizon}" # create column name for each span
          sp500[ratio_column] = sp500["Close"] / rolling_avg["Close"]

```

```

trend_column = f"Trend_{horizon}"
sp500[trend_column] = sp500.shift(1).rolling(horizon).sum()["Target"] # It
↳will calculate the past days of the Target number so if 0 0 0 1 1 then would
↳be for last 5 days there are 2 days that goes up

new_predictors += [ratio_column, trend_column]

```

[22]: sp500 # There would be lot of NaN since if they can't find the enough days
↳prior to the data, it will automatically become NaN

[22]:

	Open	High	Low	Close \
Date				
1990-01-02 00:00:00-05:00	353.399994	359.690002	351.980011	359.690002
1990-01-03 00:00:00-05:00	359.690002	360.589996	357.890015	358.760010
1990-01-04 00:00:00-05:00	358.760010	358.760010	352.890015	355.670013
1990-01-05 00:00:00-05:00	355.670013	355.670013	351.350006	352.200012
1990-01-08 00:00:00-05:00	352.200012	354.239990	350.540009	353.790009
...
2024-12-24 00:00:00-05:00	5984.629883	6040.100098	5981.439941	6040.040039
2024-12-26 00:00:00-05:00	6024.970215	6049.750000	6007.370117	6037.589844
2024-12-27 00:00:00-05:00	6006.169922	6006.169922	5932.950195	5970.839844
2024-12-30 00:00:00-05:00	5920.669922	5940.790039	5869.160156	5906.939941
2024-12-31 00:00:00-05:00	5919.740234	5929.740234	5868.859863	5881.629883

	Volume	Tomorrow	Target	Close_Ratio_2 \
Date				
1990-01-02 00:00:00-05:00	162070000	358.760010	0	NaN
1990-01-03 00:00:00-05:00	192330000	355.670013	0	0.998706
1990-01-04 00:00:00-05:00	177000000	352.200012	0	0.995675
1990-01-05 00:00:00-05:00	158530000	353.790009	1	0.995098
1990-01-08 00:00:00-05:00	140110000	349.619995	0	1.002252
...
2024-12-24 00:00:00-05:00	1757720000	6037.589844	0	1.005491
2024-12-26 00:00:00-05:00	2904530000	5970.839844	0	0.999797
2024-12-27 00:00:00-05:00	3159610000	5906.939941	0	0.994441
2024-12-30 00:00:00-05:00	3433250000	5881.629883	0	0.994620
2024-12-31 00:00:00-05:00	3128350000	NaN	0	0.997853

	Trend_2	Close_Ratio_5	Trend_5	Close_Ratio_60 \
Date				
1990-01-02 00:00:00-05:00	NaN	NaN	NaN	NaN
1990-01-03 00:00:00-05:00	NaN	NaN	NaN	NaN
1990-01-04 00:00:00-05:00	0.0	NaN	NaN	NaN
1990-01-05 00:00:00-05:00	0.0	NaN	NaN	NaN
1990-01-08 00:00:00-05:00	1.0	0.993731	NaN	NaN
...

2024-12-24 00:00:00-05:00	2.0	1.017383	3.0	1.023002
2024-12-26 00:00:00-05:00	1.0	1.011334	3.0	1.021639
2024-12-27 00:00:00-05:00	0.0	0.996689	3.0	1.009600
2024-12-30 00:00:00-05:00	0.0	0.986810	2.0	0.998213
2024-12-31 00:00:00-05:00	0.0	0.985626	1.0	0.993570

Date	Trend_60	Close_Ratio_250	Trend_250	\
1990-01-02 00:00:00-05:00	NaN	NaN	NaN	
1990-01-03 00:00:00-05:00	NaN	NaN	NaN	
1990-01-04 00:00:00-05:00	NaN	NaN	NaN	
1990-01-05 00:00:00-05:00	NaN	NaN	NaN	
1990-01-08 00:00:00-05:00	NaN	NaN	NaN	
...	
2024-12-24 00:00:00-05:00	35.0	1.115492	144.0	
2024-12-26 00:00:00-05:00	35.0	1.114008	143.0	
2024-12-27 00:00:00-05:00	34.0	1.100716	143.0	
2024-12-30 00:00:00-05:00	34.0	1.088002	143.0	
2024-12-31 00:00:00-05:00	33.0	1.082402	143.0	

Date	Close_Ratio_1000	Trend_1000
1990-01-02 00:00:00-05:00	NaN	NaN
1990-01-03 00:00:00-05:00	NaN	NaN
1990-01-04 00:00:00-05:00	NaN	NaN
1990-01-05 00:00:00-05:00	NaN	NaN
1990-01-08 00:00:00-05:00	NaN	NaN
...
2024-12-24 00:00:00-05:00	1.337188	531.0
2024-12-26 00:00:00-05:00	1.335962	530.0
2024-12-27 00:00:00-05:00	1.320543	529.0
2024-12-30 00:00:00-05:00	1.305803	528.0
2024-12-31 00:00:00-05:00	1.299617	527.0

[8817 rows x 17 columns]

```
[23]: # Fixing the NaN issues
sp500 = sp500.dropna()
```

1.0.5 Improving Models

```
[24]: model = RandomForestClassifier(n_estimators=200, min_samples_split=50,
↳ random_state=1)
```

```
[25]: # Similar as above but instead find the probability of the prediction
def predict(train, test, predictors, model):
    model.fit(train[predictors], train["Target"])
```



```

    preds = model.predict_proba(test[predictors])[:,1] #predict_proba.␣
    ↪ Probability the goes up
    preds[preds >= 0.6] = 1
    preds[preds < 0.6] = 0
    # Increase the chance that goes up

    preds = pd.Series(preds, index=test.index, name="Predictions") #Named the␣
    ↪ series as Predictions
    combined = pd.concat([test["Target"], preds], axis = 1)
    return combined

```

```

[26]: # Backtest again
      predictions = backtest(sp500, model, new_predictors)

```

```

[27]: predictions["Predictions"].value_counts()
      #Why there's lot of 0's?
      #A: since we change the threshold to make it more confidence

```

```

[27]: Predictions
      0.0    4463
      1.0     853
      Name: count, dtype: int64

```

```

[29]: precision_score(predictions["Target"], predictions["Predictions"])

```

```

[29]: 0.5756154747948418

```

```

[46]: from sklearn.metrics import mean_absolute_error, mean_squared_error, f1_score

      mae_test = mean_absolute_error(test["Target"],preds)
      f1_test = f1_score(test["Target"], preds)

      mae_new = mean_absolute_error(predictions["Target"], predictions["Predictions"])
      f1_new = f1_score(predictions["Target"], predictions["Predictions"])

```

```

[50]: print(f"Mean Absolute Error: {mae_test}, F1 Score: {f1_test}")

```

```

Mean Absolute Error: 0.45, F1 Score: 0.48275862068965514

```

```

[49]: print(f"Mean Absolute Error: {mae_new}, F1 Score: {f1_new}")

```

```

Mean Absolute Error: 0.5212565838976674, F1 Score: 0.2616573407940314

```

1.1 Conclusion

Adding more predictors does not always lead to better performance. In this case, the added predictors likely introduced noise or complexity, outweighing any potential benefits.

[]: